



Charles W. Davidson College of Engineering
Department of Computer Engineering

**Real-Time Embedded System
Co-Design
CMPE 146 Section 1
Fall 2024**




Memory


- SRAM stands for Static Random Access Memory
- Volatile
 - When power is gone, data are gone
- Why static?
 - Data are maintained without any kind of extra operations
 - As long as power to the device is maintained

- DRAM stands for Dynamic Random-Access Memory
- Volatile
 - When power is gone, data are gone
- Why dynamic?
 - Data are maintained with periodic refresh operations
 - Just maintaining power to the device is not sufficient
- Comparison to SRAM
 - DRAM cell is smaller
 - Less expensive per bit
 - SRAM is faster
 - DRAM requires more peripheral circuitry
- Hard to integrate with logic circuitry on chip
 - Fabrication process is different
 - Usually integrated off-chip

- ROM stands for Read-Only Memory
- Non-volatile
 - When power is gone, stored data are still retained
- Some types are programmable and non-erasable
 - PROM (Programmable Read-Only Memory)
 - Once programmed, memory contents cannot be changed
- Some types are programmable and erasable
 - EPROM (Erasable Programmable Read-Only Memory)
 - Erasing and programming are done off board
 - Erase with ultra-violet light (in 15-20 minutes)
 - Program with high voltage (12 V)
 - EEPROM (Electrically Erasable Programmable Read-Only Memory)
 - Basically an older version of flash memory
 - Erasing and programming are done on board
 - Erasing is much slower and needs higher voltage (20 V)
 - Lower cell density




- 6 transistors




BL: Bit Line, carries data

WL: Word Line, selects cells

- On write
 - Set BLs to new data value
 - Raise WL to High
 - Cell is set to new state




Stable Configurations




- On read
 - Set BLs High (precharge)
 - Set WL High and float BLs
 - One of the BLs goes Low
 - Sense the difference of BLs

- 1 transistor, 1 capacitor
- On write
 - Set BL High or Low
 - Raise WL to High
 - Cell is set to new state




- On Read
 - Precharge BL to half of V_{dd}
 - Set WL High and float BL
 - Capacitor shares charge with BL, causing a change of voltage in BL
 - Sense BL voltage
 - Higher than half of V_{dd} : High is read
 - Lower than half of V_{dd} : Low is read
 - Restore charge state in capacitor (recharge or drain)
 - Execute a write operation of the same data just read

- 1 transistor with floating gate (FG)
 - Electrons at FG means transistor cannot be conducting
 - Neutralize the effect of an application of voltage at gate
- NOR flash (used for program execution) operations:
- On Erase
 - Float BL, set WL Low and apply 12 V to SL
 - Effect: Clear all electrons at FG
 - Cell stores a 1 (transistor can be turned ON)
- On Write/Program
 - Apply ~12 V to WL and set SL to Low
 - Store 0: Apply ~7 V to BL → Electrons at FG
 - Otherwise, set BL to Low → No change at FG
- On Read
 - Apply ~5V to WL, ~ 1V to BL and set SL to Low
 - Stored 0 (electrons at FG): Transistor is turned off, no current from D to S
 - Stored 1 (no electrons at FG): Transistor is turned on, some current from D to S
 - Presence of current flow in BL means a stored 1





G: Gate, D: Drain, S: Source



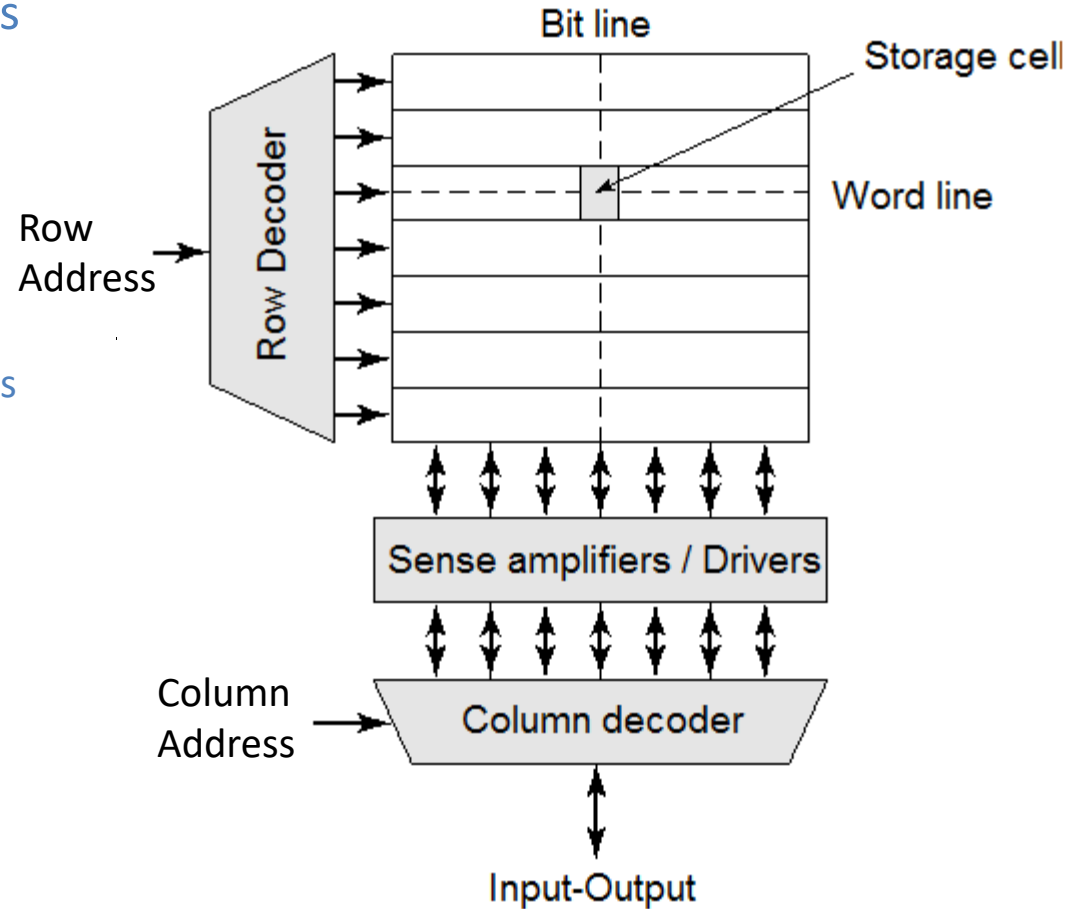
SL: Source Line

- Implementation with at most 1 transistor
 - Data “written” during chip design
- Cell containing a transistor stores a 1
- Cell containing no transistor stores a 0



- On Read
 - Set WL High
 - BL becomes High or Low
 - With BL attached to a pull-down resistor

- Cells are arranged in a 2-D array
- Word line can be various lengths
 - 8, 16, 32, 64 or 128 bits
- Sense amplifiers
 - Detect changes on bit lines
 - Differential type for SRAM
 - Memory cell has two bit lines
- Achieves high density of cells




A controller manages the operations on one or more memory cell arrays. It has a number of functions.

- Interfaces with the rest of the system
 - For example, connects to the AHB (Advanced High-Performance Bus)
- Provides various operation modes
 - Asynchronous or synchronous (to the system clock)
 - Burst data transfer
 - Low power mode (for example, only retains data)
- Control
 - Command decoding
 - Configuration (from processor)
 - Number of wait states
 - Erase protection bits (for flash memory)
 - Data buffering
 - Provides caching scheme to improve access performance



- Timing
 - Read, write, erase (for flash memory)
 - May need to insert wait states with high-frequency bus clocks
 - Refresh (for DRAM)
 - DDR (Double Data Rate) operations (for DRAM)
 - Uses both the falling and rising edges of clock
 - Command pipelining
- Address decoding
 - Maps input address to into bank, row, column
- Error detection and correction
 - Uses parity (single bit) or ECC (Error Correction Code, multiple bits)
- Interfaces with target memory (SRAM, DRAM, flash, ROM)
 - Row and column address decoders
 - Sense amplifiers and drivers (on bit lines)
 - Higher voltages (for flash memory)

Memory Interleaving


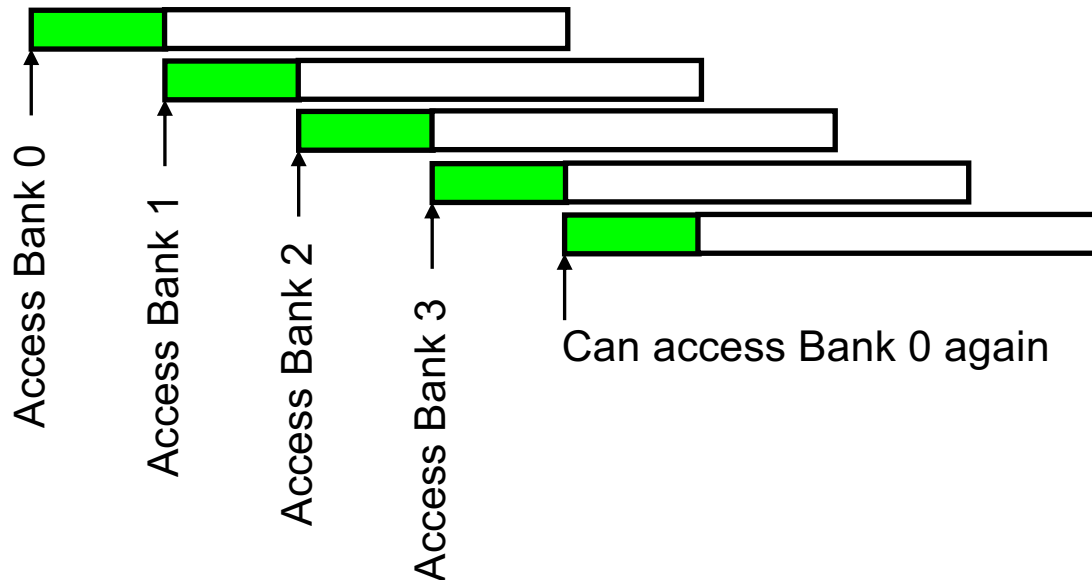
- Put a line of words in different memory banks
 - To achieve better performance



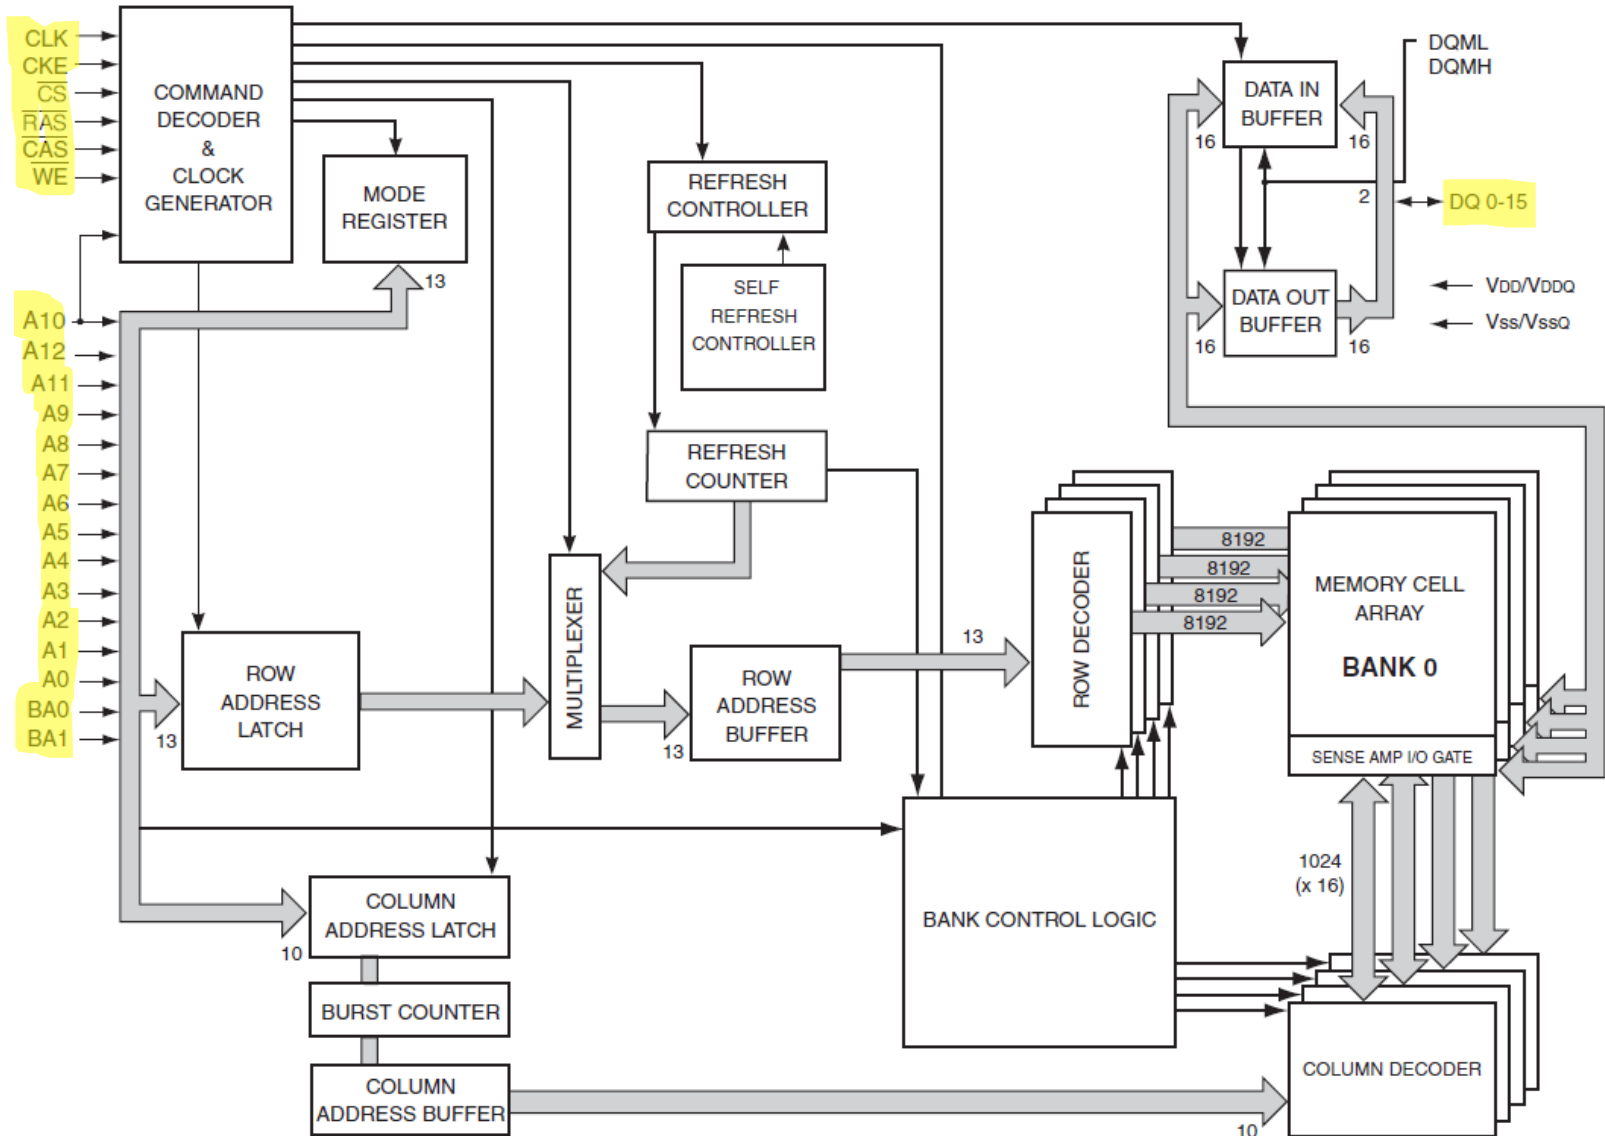
- Access pattern without interleaving:




- Access pattern with 4-way interleaving:



- Integrated Silicon Solution IS42/45X 512-Mb Synchronous DRAM






Computing performance is usually limited by memory latency and bandwidth

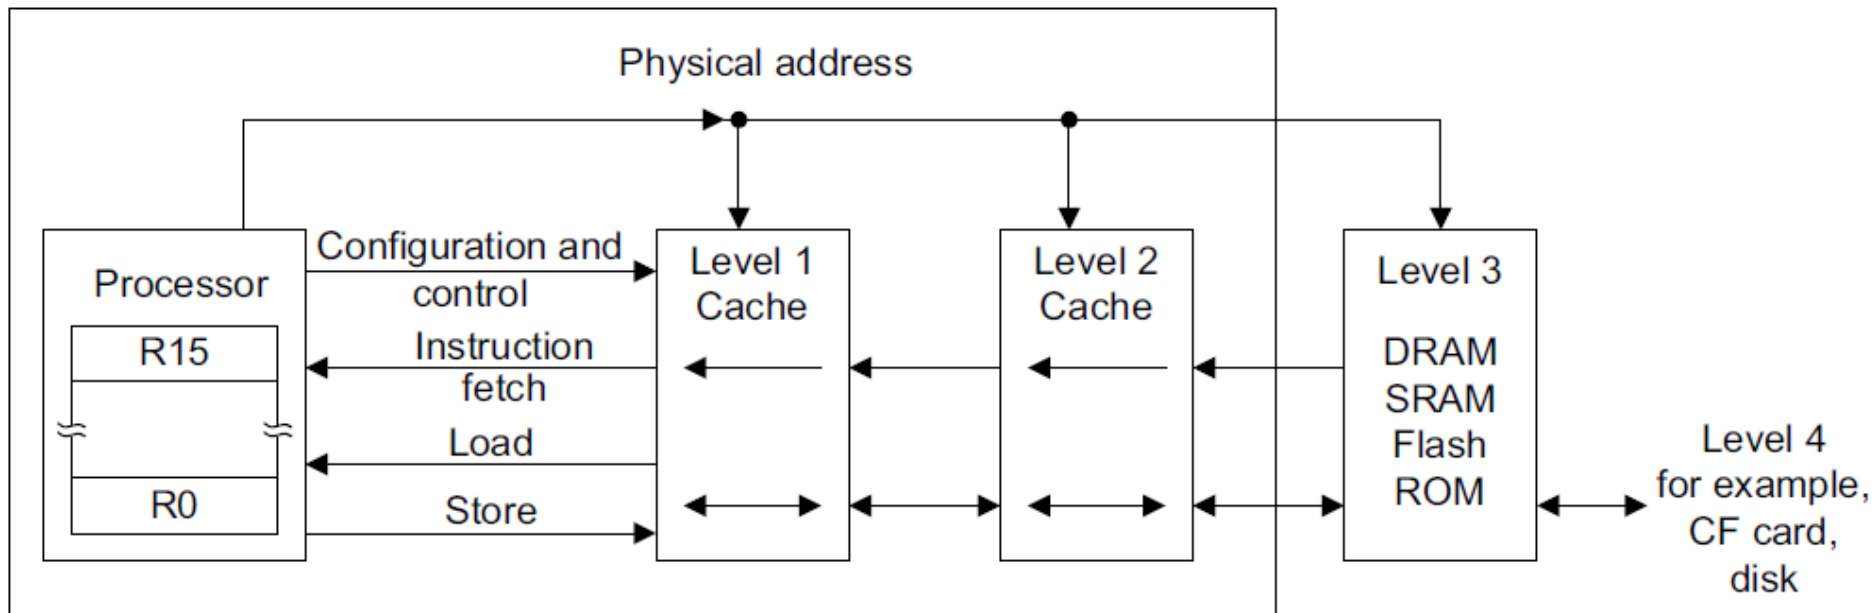
- Latency
 - Access time: Time between request made and data arrives
 - Affected by cycle time: Minimum time between requests to memory
- Bandwidth
 - Data transferred per unit time
- Modern embedded processors can run at hundreds of MHz or higher
 - Need very fast memory to avoid stalling the processor execution
 - Normal main memory with at least tens of ns access time is not adequate

- A very effective solution to much narrow the gap between processor and memory speeds
- Create a very fast and small memory (cache) sitting between the processor and the main memory
 - Cache operates at (or very close to) the processor speed




- How it works?
 - When processor reads, if the data is not in the cache, it reads a block of data (cache line) from memory including the desired data into the cache [**Cache Miss**]
 - Returns the desired data to processor
 - Next time when processor reads, if the data is in the cache, data is returned to processor immediately [**Cache Hit**]
 - When processor writes, it can write to cache first immediately, without going all the way to main memory

- The closer to processor core, the faster the memory is
 - Amount of availability is also less
 - More expensive per bit
- Gives the illusion of a large, fast memory being presented to the processor
- Goal: To reduce the average memory access time



Why Cache Works?

- Most application executions exhibit two properties: Temporal and spatial localities
- Temporal locality
 - If a location is referenced, it is likely to be referenced again in the near future
 - For examples, small program loops, array index variable in a loop, table pointer, etc.
 - Cache “remembers” the contents of recently accessed locations
- Spatial locality
 - If a location is referenced, it is likely that locations near it will be referenced in the near future
 - For example, structured data such as arrays, matrices, queues, etc.
 - Cache also fetches data around recently accessed locations



```

sum = 0;
for (int i=0; i<N; i++)
    sum += data[i];
    
```

- Need to manage potentially two copies of same data
 - Only modified data are problematic
 - Management complexity increases quickly with more processors in system
- DMA
 - Controller typically has no access to cache, which is tightly coupled with the processor, not attached to the system bus
 - Possible solutions
 - Introduce circuitry to detect potentially problematic conditions on bus
 - Make application program be aware of potential conflicts