



Charles W. Davidson College of Engineering
Department of Computer Engineering

**Real-Time Embedded System
Co-Design
CMPE 146 Section 1
Fall 2024**



Task Scheduler

- Let's loosely define a task as a group of program instructions designed to achieve a specific objective
- Tasks are typically determined and well defined when the system is being built in the design phase
- In embedded applications, quite often, execution of a task is triggered by some event
 - Typically it lasts for a finite duration before the same event occurs again
- A task generally can be organized as a function
 - In most embedded systems, there is only one program to execute
 - All the tasks are contained in the same program
 - Unlike in conventional computer systems where tasks can be organized in multiple executable program files
 - Note that a single function can be used to perform multiple tasks
 - For example:

```
process_sensor_input(Sensor_A);  
process_sensor_input(Sensor_B);  
process_sensor_input(Sensor_C);
```

- Effective way to build a software system with the modular concept
 - Similar to hardware design that has different functional modules
- Clear division of system development effort and functionality
 - Allows team of developers to work in parallel
- System design can be understood more easily
 - Make future enhancements easier
 - Reduces future maintenance cost
 - Embedded products can last for many years

We generally have three types of task in embedded applications

- Periodic tasks
 - Time between two consecutive executions is fixed
 - The time (period) is known
 - For example, a weather station gets a temperature reading once every minute
- Sporadic tasks
 - Time between two consecutive executions may differ widely
 - The time is totally random
 - Can be arbitrarily short or long
 - For example, task triggered by a touchpad input device of a game console
 - The device may be untouched for days , then it is being used in rapid fashion when the user is playing game
- Aperiodic tasks
 - Time between two consecutive executions may follow a known probability distribution function
 - For example, the data packet arrival rate at a communication port may follow a Poisson distribution

For simple embedded applications, a simple scheduling mechanism is probably adequate

- Round Robin
- Round Robin with Interrupts
- Queue-based

- Well-known principle that is widely adopted in many disciplines
 - A very simple scheduling algorithm that can be easily implemented
- All tasks are treated equally
 - Every task takes its turn sequentially to be run
 - A while loop (super loop) in main() calls the task functions sequentially
 - Each function checks for particular event(s). If no event occurs, it just exits immediately.

```
while (true)
{
    task1();
    task2();
    ⋮
    taskN();
}
```

- Suitable for simple embedded systems that can tolerate potentially slow response time
 - A task's worst-case response time depends on sum of execution time of all other tasks

- Based on the simple round-robin scheme, add interrupt handling to provide fast response (small delay) for certain events
 - Tasks are in the foreground (or at higher level) where most normal processor execution takes place
 - Interrupt handlings are in the background (or at lower level)
- All tasks are treated equally at the higher level
 - Interrupt handlers pre-empt the round-robin loop to process the interrupt events

```
while (true)
{
    task1();
    task2();
    ⋮
    taskN();
}
```

```
void isr1()
{
    ⋮
}
```

```
void isr2()
{
    ⋮
}
```

....

```
void isrM()
{
    ⋮
}
```

- Task and ISR can work together to accomplish a more complex task that allows for a longer response time
 - ISR provides immediate response and task responses as more signals or events are collected
 - Use shared variables or data structures to communicate

- One example: Communication monitor handles commands from UART
 - When UART receives a character in the command string, it creates an interrupt
 - ISR quickly acknowledges UART controller and puts the received character in a shared buffer
 - A task keeps checking the buffer to see if a terminating newline is received
 - If newline received, the task handles the command and respond accordingly
 - Otherwise, it keeps checking continually
 - Without using interrupt, as in the simple round-robin scheme, some incoming characters may be lost
 - Characters can come in at such a high speed that the round-robin loop is not quick enough to invoke the handling task
- As in simple round-robin scheme, some tasks' worst-case response time remains the same, but urgent events can be handled very quickly
 - Overall system performance is better than the simple round-robin scheme
 - Extra circuitry is needed to handle the interrupt signals

- We can add a queue to further improve the round-robin-with-interrupts scheme
- An ISR can invoke an upper-level task by placing the task's address in a FIFO (First-In-First-Out) queue when the task has something to do

```
while (true)
{
    if (queue_empty())
        continue;
    function = get_first_entry_from_queue();
    call function;
}
```

Task queue



- In the previous communication monitor example, the ISR will place the handling task's address in the queue when a newline is received
- Reduces many unnecessary checking in the super loop when tasks are not ready to execute yet
 - No need to go through each task in the loop in each iteration
 - Can use low-power mode to reduce power consumption
 - Foreground super loop is only active when there is something to do

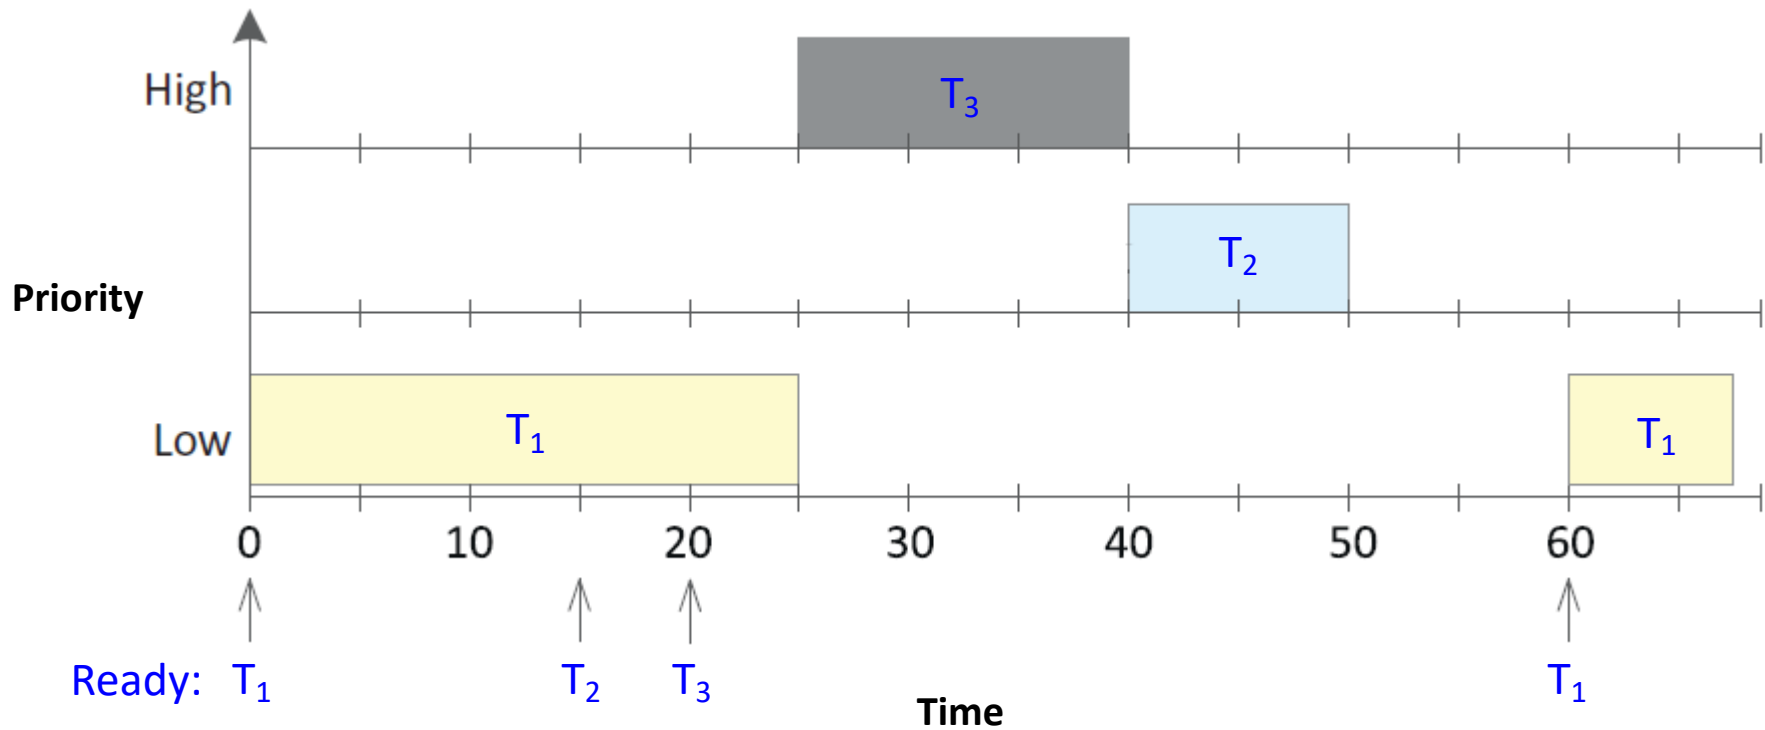
- Better system response can be achieved by adding a priority attribute to a task
- With additional complexity, application development would be much easier with an OS
- In POSIX (Portable Operation System Interface), threads (tasks) are schedulable entities
- The OS scheduler selects one thread to execute on each processor at some point in time
- POSIX thread is associated with a numeric value, the scheduling priority
 - Used by the scheduler to determine when and how to schedule the thread

- Three principles are used in priority-based scheduling in POSIX
 - Precedence principle
 - When no thread is running, and multiple threads are ready to be scheduled for execution, a thread with a higher priority is always scheduled prior to threads with a lower priority
 - Preemption principle
 - When a thread is currently running, and another thread with a higher priority becomes ready for execution, the higher-priority thread preempts the execution of the current thread
 - Fairness principle
 - When multiple threads, with the same priority, compete for use of the processing resources, the system's scheduling behavior is regulated by the scheduling policies

Precedence Principle

- The precedence principle is the basis of priority-based scheduling
- A non-preemptive scheduling approach illustrates the principle:

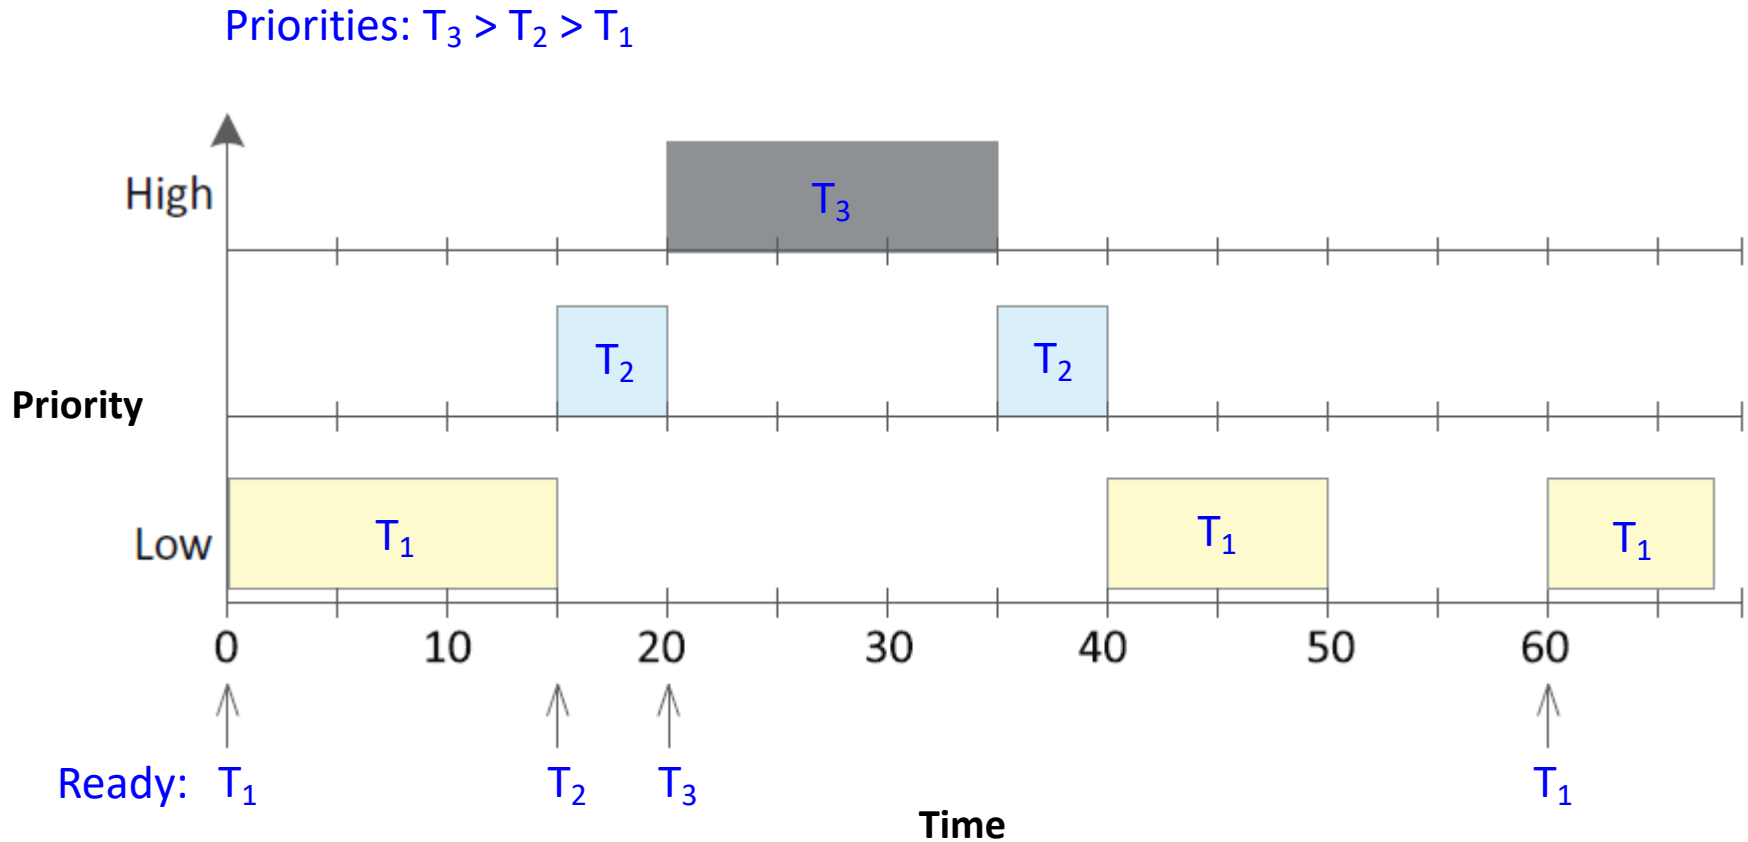
Priorities: $T_3 > T_2 > T_1$



- T_3 gets to execute before T_2 even though T_3 was ready later

Preemption Principle

- A preemptive scheduling approach illustrates the principle:

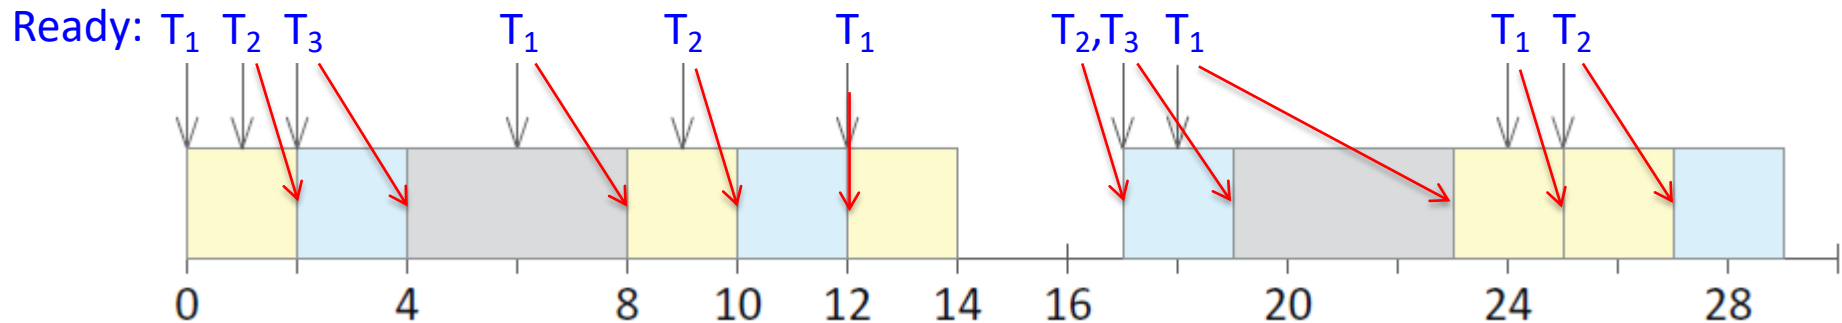
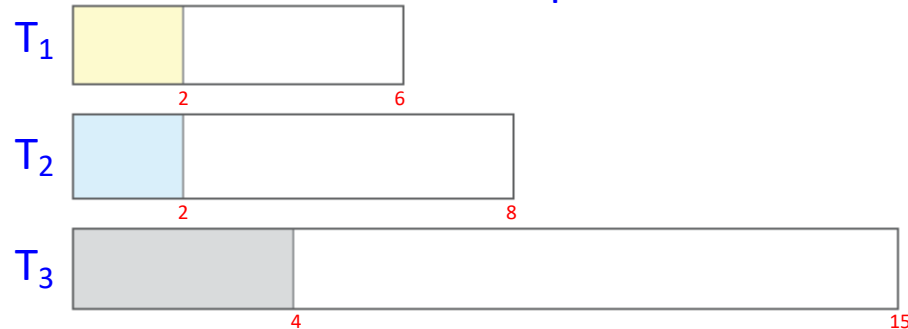


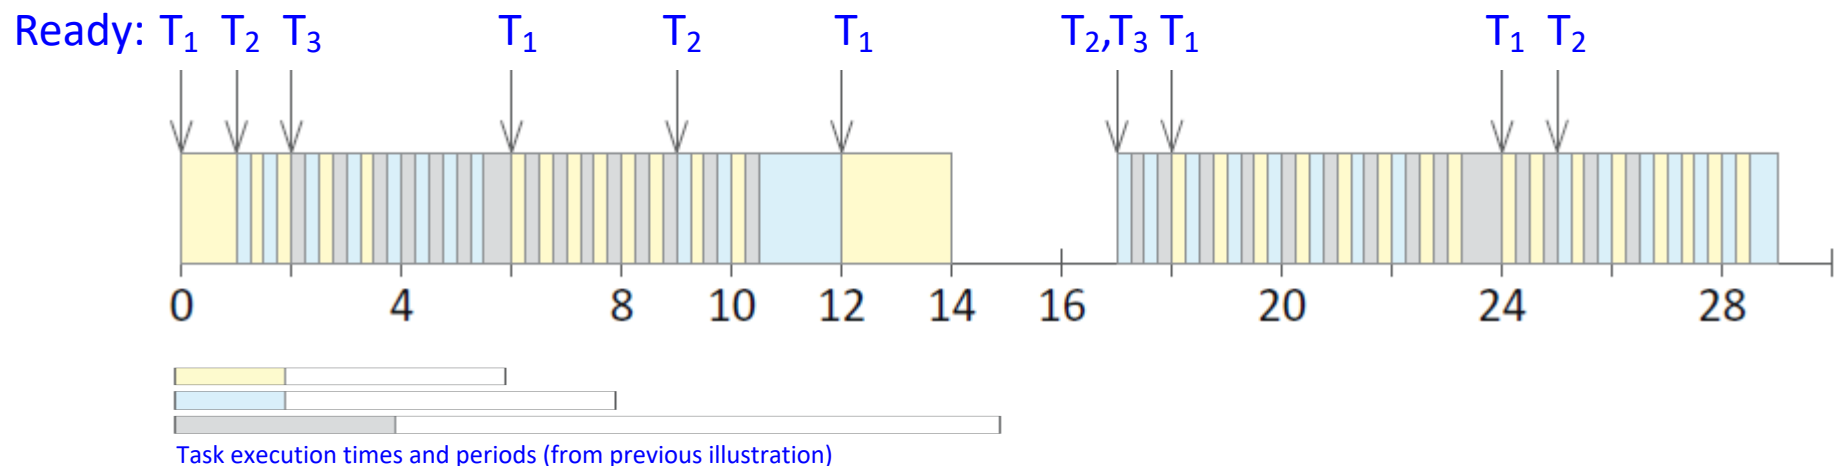
- We use scheduling policies to define the fairness rules for tasks with the same priority
- Two basic scheduling policies
 - FIFO (First-In-First-Out)
 - Round Robin

FIFO Scheduling Policy

- Applied to scheduling tasks with the same priority
- A task cannot be preempted by any other tasks with later entry time in the task-ready list
 - The task that entered the task-ready list first will run to its completion before other tasks are scheduled to run

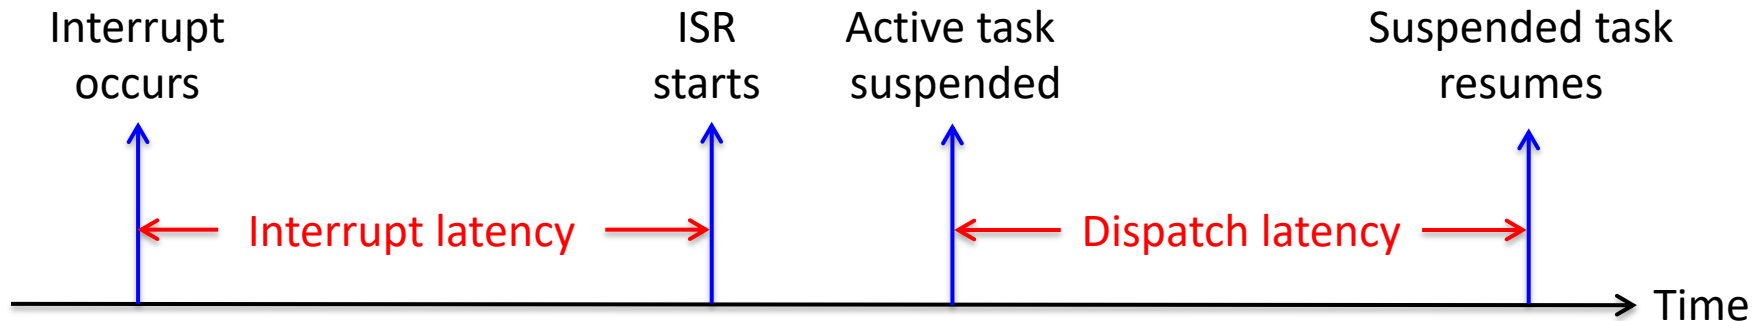
Task execution times and periods

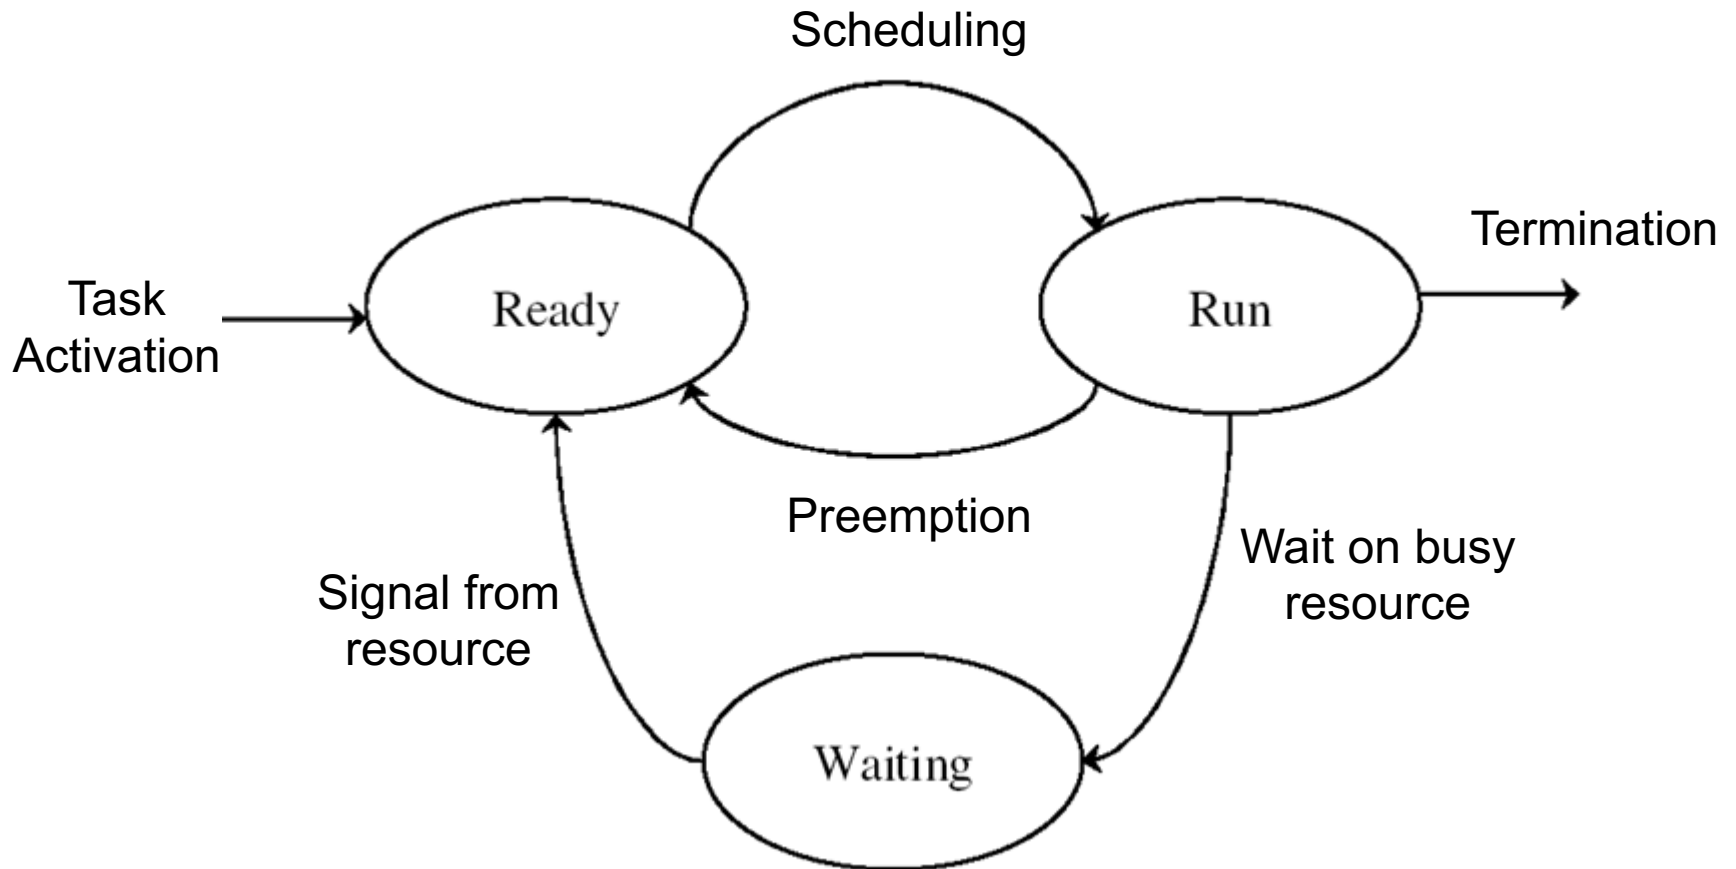




- A very important feature of a RTOS is to be able to respond immediately to an event, i.e., to run the handling task as soon as possible
- Urgency or importance of a task is usually indicated by its priority
 - Tasks need to run sooner are assigned with higher priorities
- As a result, the RTOS must support
 - Priority-based scheduler
 - Preemption of tasks
- To ensure real-time performance, a task currently running needs be preempted if a higher-priority task becomes ready to run
- Providing a preemptive, priority-based scheduler only guarantees soft real-time functionality (best-effort approach)
 - Hard real-time systems must further guarantee that real-time tasks will be serviced in accord with their deadline requirements
 - Making such guarantees may require additional scheduling features

- When an event occurs, the system must respond to and service it as quickly as possible
- Two types of latencies affect the real-time performance under RTOS
 - Interrupt latency
 - Amount of time from the arrival of an interrupt to the start of ISR
 - Dispatch latency
 - Amount of time required for the RTOS to suspend one task and resume another





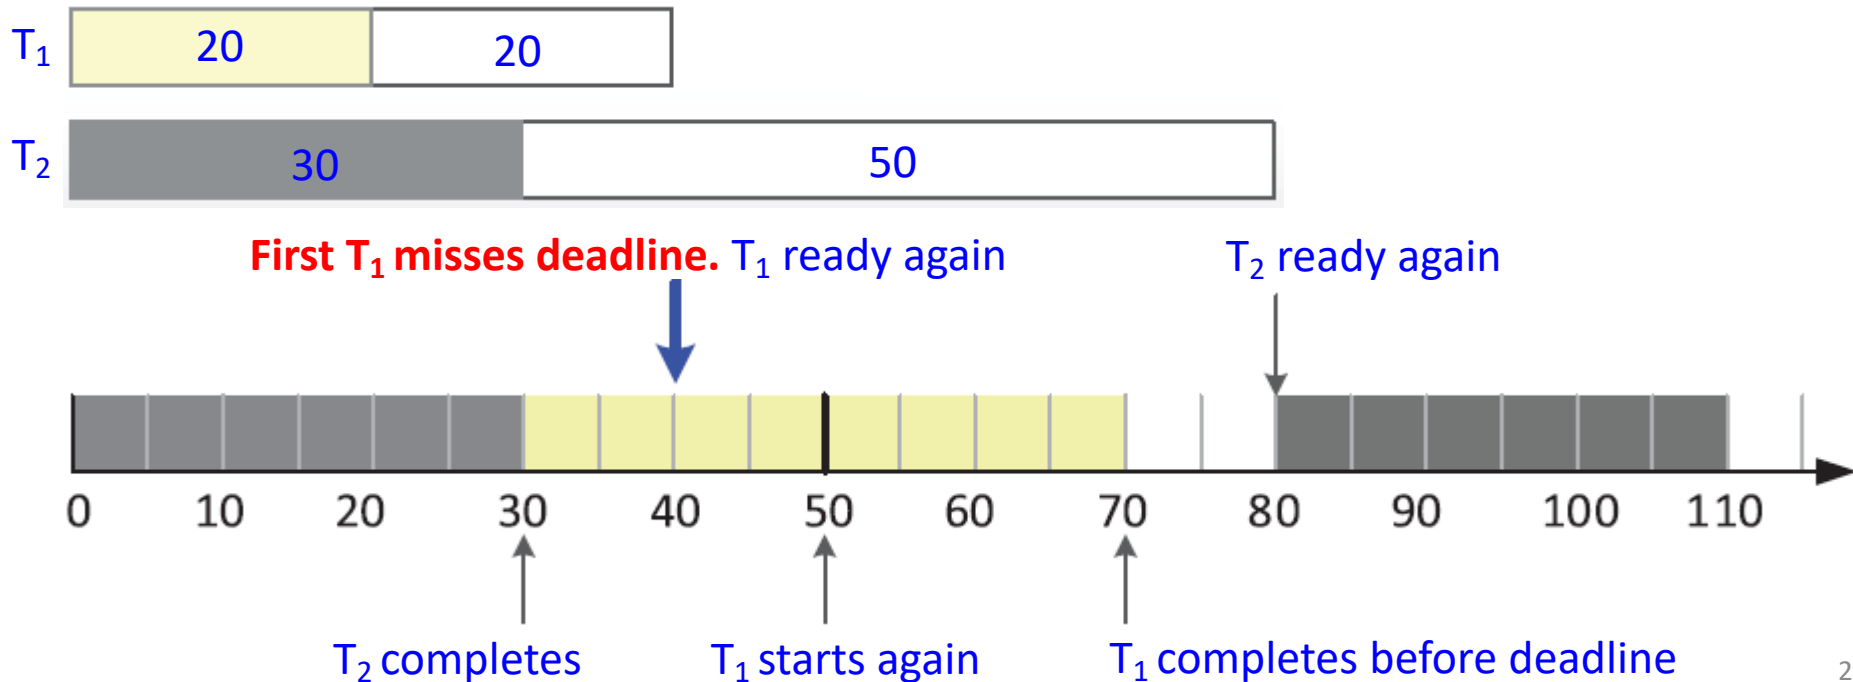
- Attributes of tasks to be scheduled
 - Period (p)
 - Deadline (d)
 - Fixed processing time (t)
- Relationship among the parameters: $0 < t \leq d \leq p$
 - Scheduler uses these relationships and assign priorities according to the periods (p) or deadlines (d)
- Three common scheduling algorithms
 - Rate-Monotonic
 - Earliest-Deadline-First
 - Deadline-Monotonic

- Schedule periodic tasks using a static priority policy with preemption
- If a lower-priority task is running and a higher-priority task becomes available to run, the scheduler will preempt the lower-priority task
- Each task is assigned a priority inversely based on its period
 - The shorter the period, the higher the rate and thus the priority
 - Priority increases with rate monotonically
 - Based on the rationale that tasks need to run more often have higher priority in order to avoid missing the deadline
- If the sum of processor utilization of tasks is less than a bound given by $k(2^{1/k} - 1)$, the task set is definitely schedulable
 - where k =number of tasks and processor utilization = t_i/p_i
 - Note that the bound is rather pessimistic; exceeding it may still be schedulable
 - When k approaches infinity, utilization sum is 69%
 - In practice, the bound is about 88% on average
- Rate-monotonic scheduling is considered optimal
 - If a set of tasks cannot be scheduled by this algorithm, it cannot be scheduled by any other algorithm that assigns static priorities

Period (p)
Deadline (d)
Fixed processing time (t)

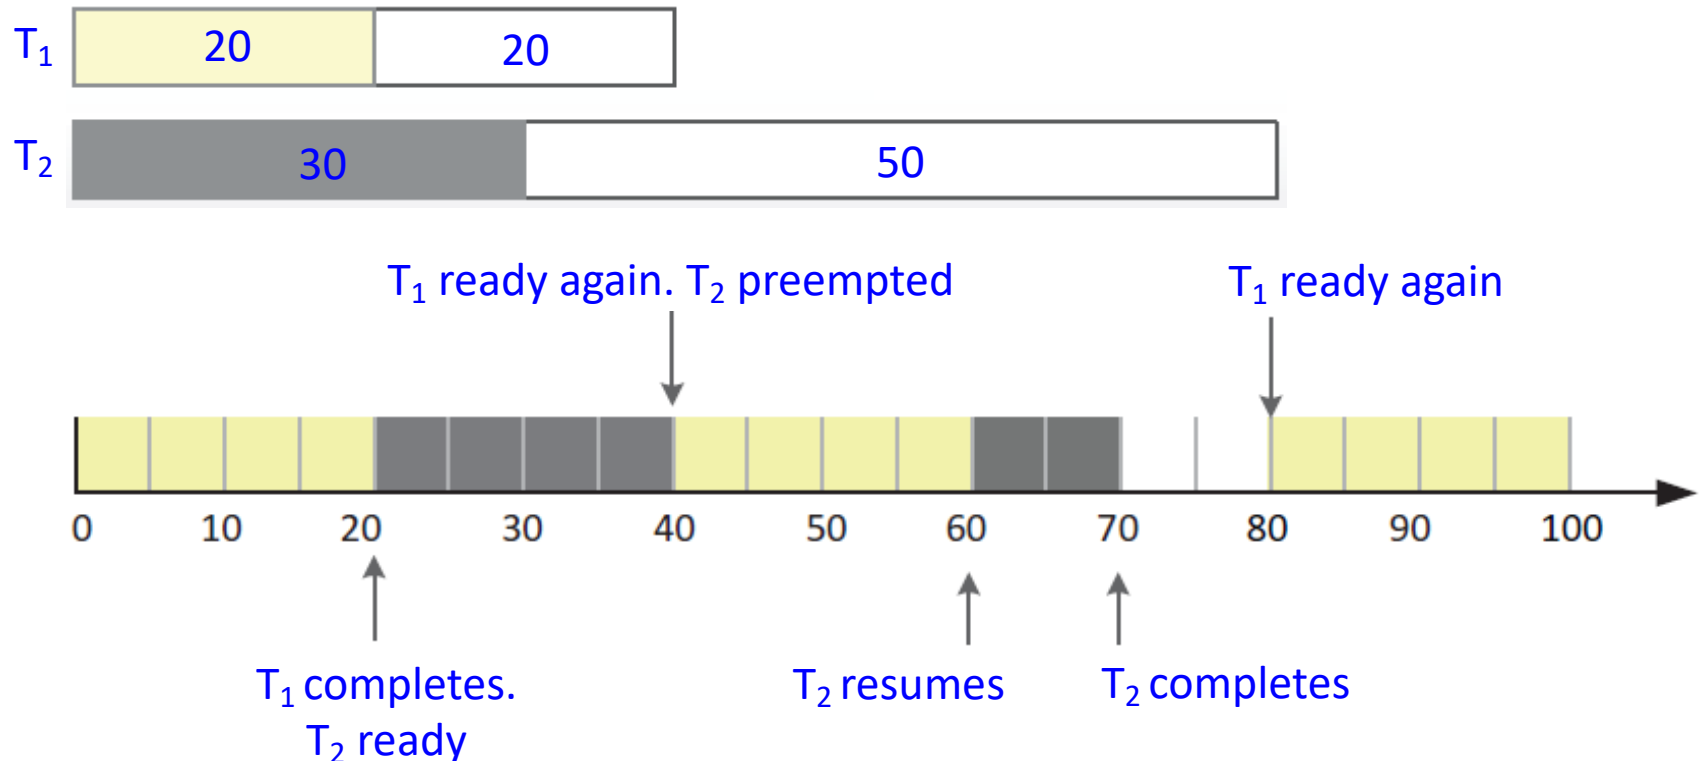
- Example: Tasks 1 and 2 with the following periods and processing times
 - T_1 : $p_1 = d_1 = 40$, $t_1 = 20$
 - T_2 : $p_2 = d_2 = 80$, $t_2 = 30$
- Both tasks are ready at Time 0
- Scenario 1: Set $\text{Priority}(T_1) < \text{Priority}(T_2)$ — not rate-monotonic
- Result: T_1 misses deadline

Period (p)
Deadline (d)
Fixed processing time (t)



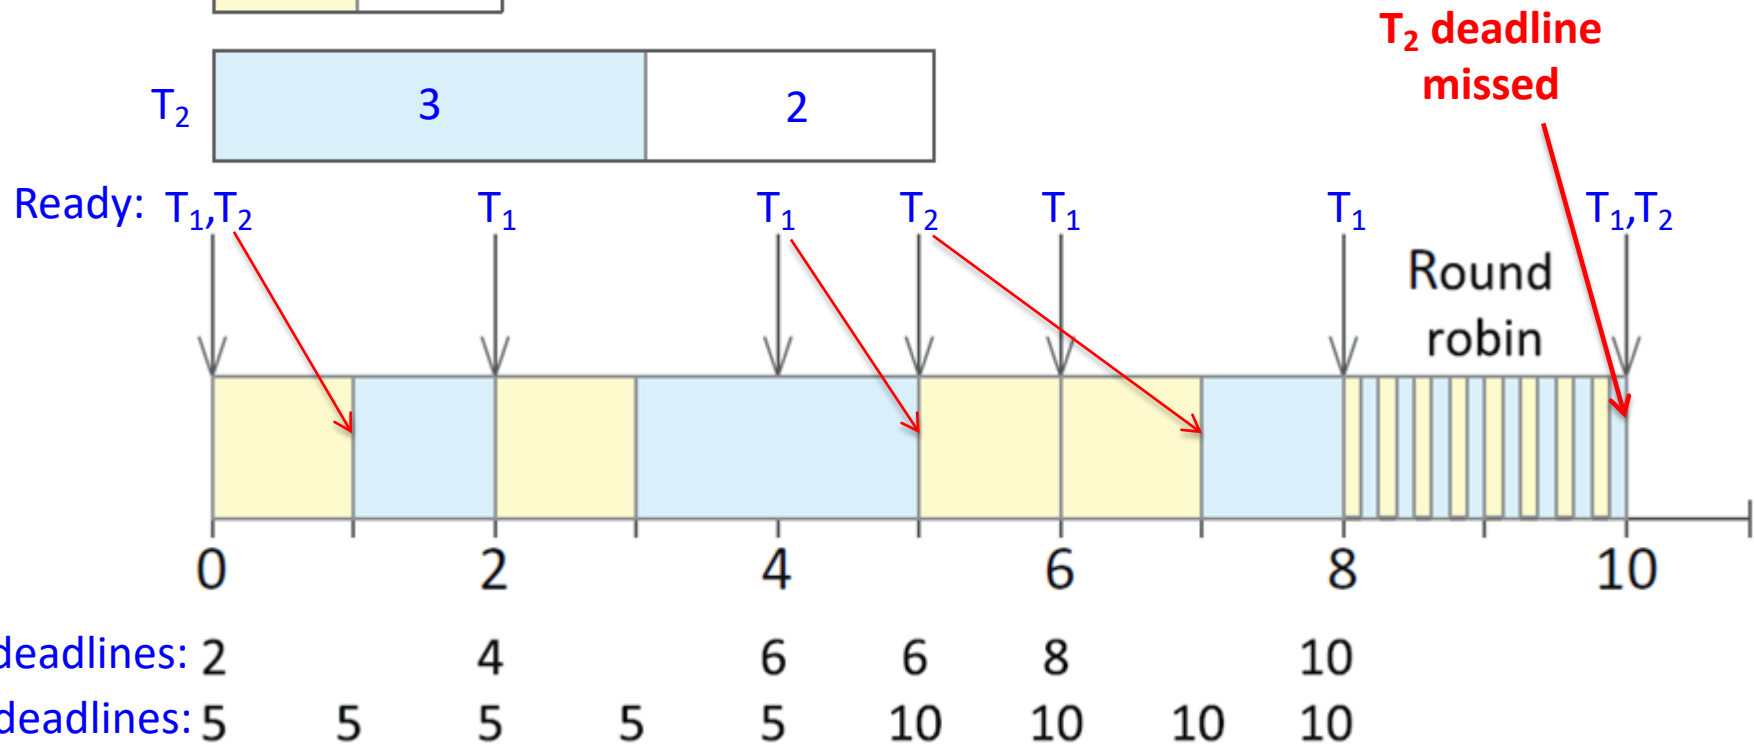
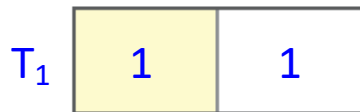
- Scenario 2: Set $\text{Priority}(T_1) > \text{Priority}(T_2)$ — rate-monotonic
- Test bound: $k(2^{1/k} - 1) = 2(2^{1/2} - 1) = 0.83$
- Sum of processor utilizations = $t_1/p_1 + t_2/p_2 = 20/40 + 30/80 = 0.88$
- Result: No missing deadlines

Period (p)
Deadline (d)
Fixed processing time (t)



- A dynamic-priority scheduling algorithm
- Requires that a task with an earlier deadline has a higher priority
- Example: Tasks 1 and 2 with the following periods and processing times
 - T_1 : $p_1 = d_1 = 2$, $t_1 = 1$
 - T_2 : $p_2 = d_2 = 5$, $t_2 = 3$

Period (p)
Deadline (d)
Fixed processing time (t)



- A static-priority scheduling algorithm
- Priorities are assigned to tasks according to their deadlines known at design time
 - A higher priority is assigned to a task with an earlier deadline
 - Priority increases with $1/\text{deadline}$ monotonically
 - Unlike the EDF approach, task priorities are fixed
- Example: Tasks 1, 2 and 3 with the following deadlines and processing times
 - T_1 : $p_1 = 60$, $t_1 = 25$, $d_1 = 50$
 - T_2 : $p_2 = 60$, $t_2 = 10$, $d_2 = 40$
 - T_3 : $p_3 = 60$, $t_3 = 15$, $d_3 = 60$
- According to their deadlines, $d_2 < d_1 < d_3$
 - Task priorities: $T_2 > T_1 > T_3$

Period (p)
Deadline (d)
Fixed processing time (t)

- Task priorities: $T_2 > T_1 > T_3$

Processing times and deadlines

